



How Far Can You Trust Your Deep Neural Networks

Feat. TensorFlow

Sungjoon Choi

Back to the Basic

I made two mistakes

PyCon KR 2017 Call for Proposals

Email address *

sungjoon.choi@cplab.snu.ac.kr

이름 (Name) *

Sungjoon Choi

세션 제목 (Title of your talk) *

How far can you trust your Deep Neural Networks? (feat. TensorFlow)

세션 카테고리 (Which category best describes your talk?) *

- Best Practices & Patterns
- Python Core (language, stdlib, etc.)
- Python Libraries
- Science / Data
- Community
- Service and Web Development
- Environment and DevOps
- Other:

세션 난이도 (Level of difficulty) *

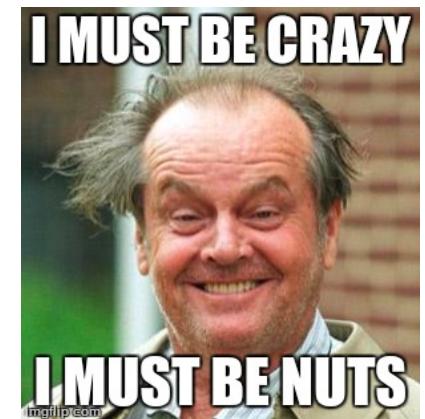
- 초급 (Beginner)
- 중급 (Intermediate)
- 고급 (Experienced)

세션 길이 (How long would you talk about the topic?) *

- 25분 (25mins)
- 40분 (40mins)

언어 (Language) *

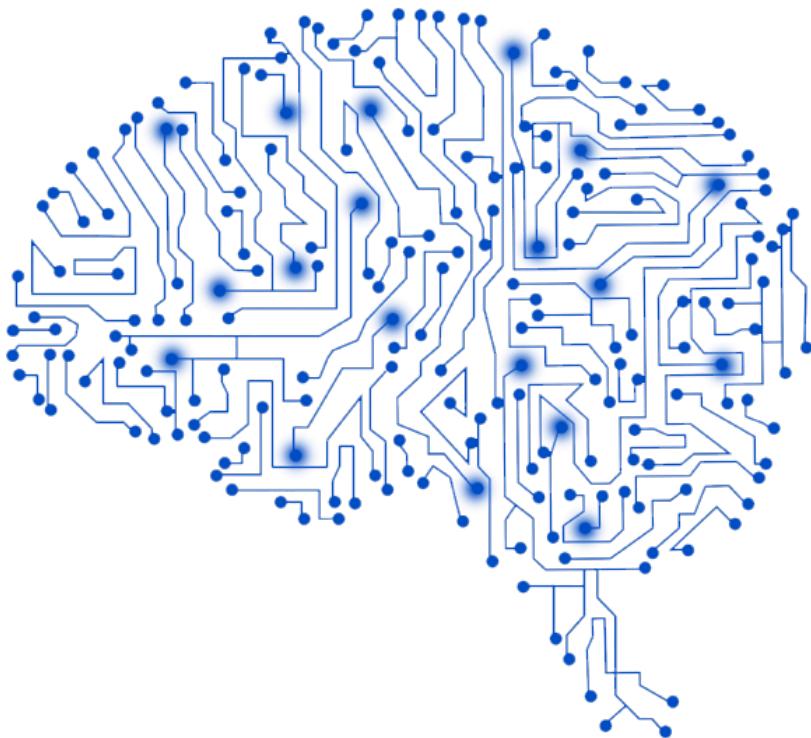
- 영어 (English)
- 한국어 (Korean)

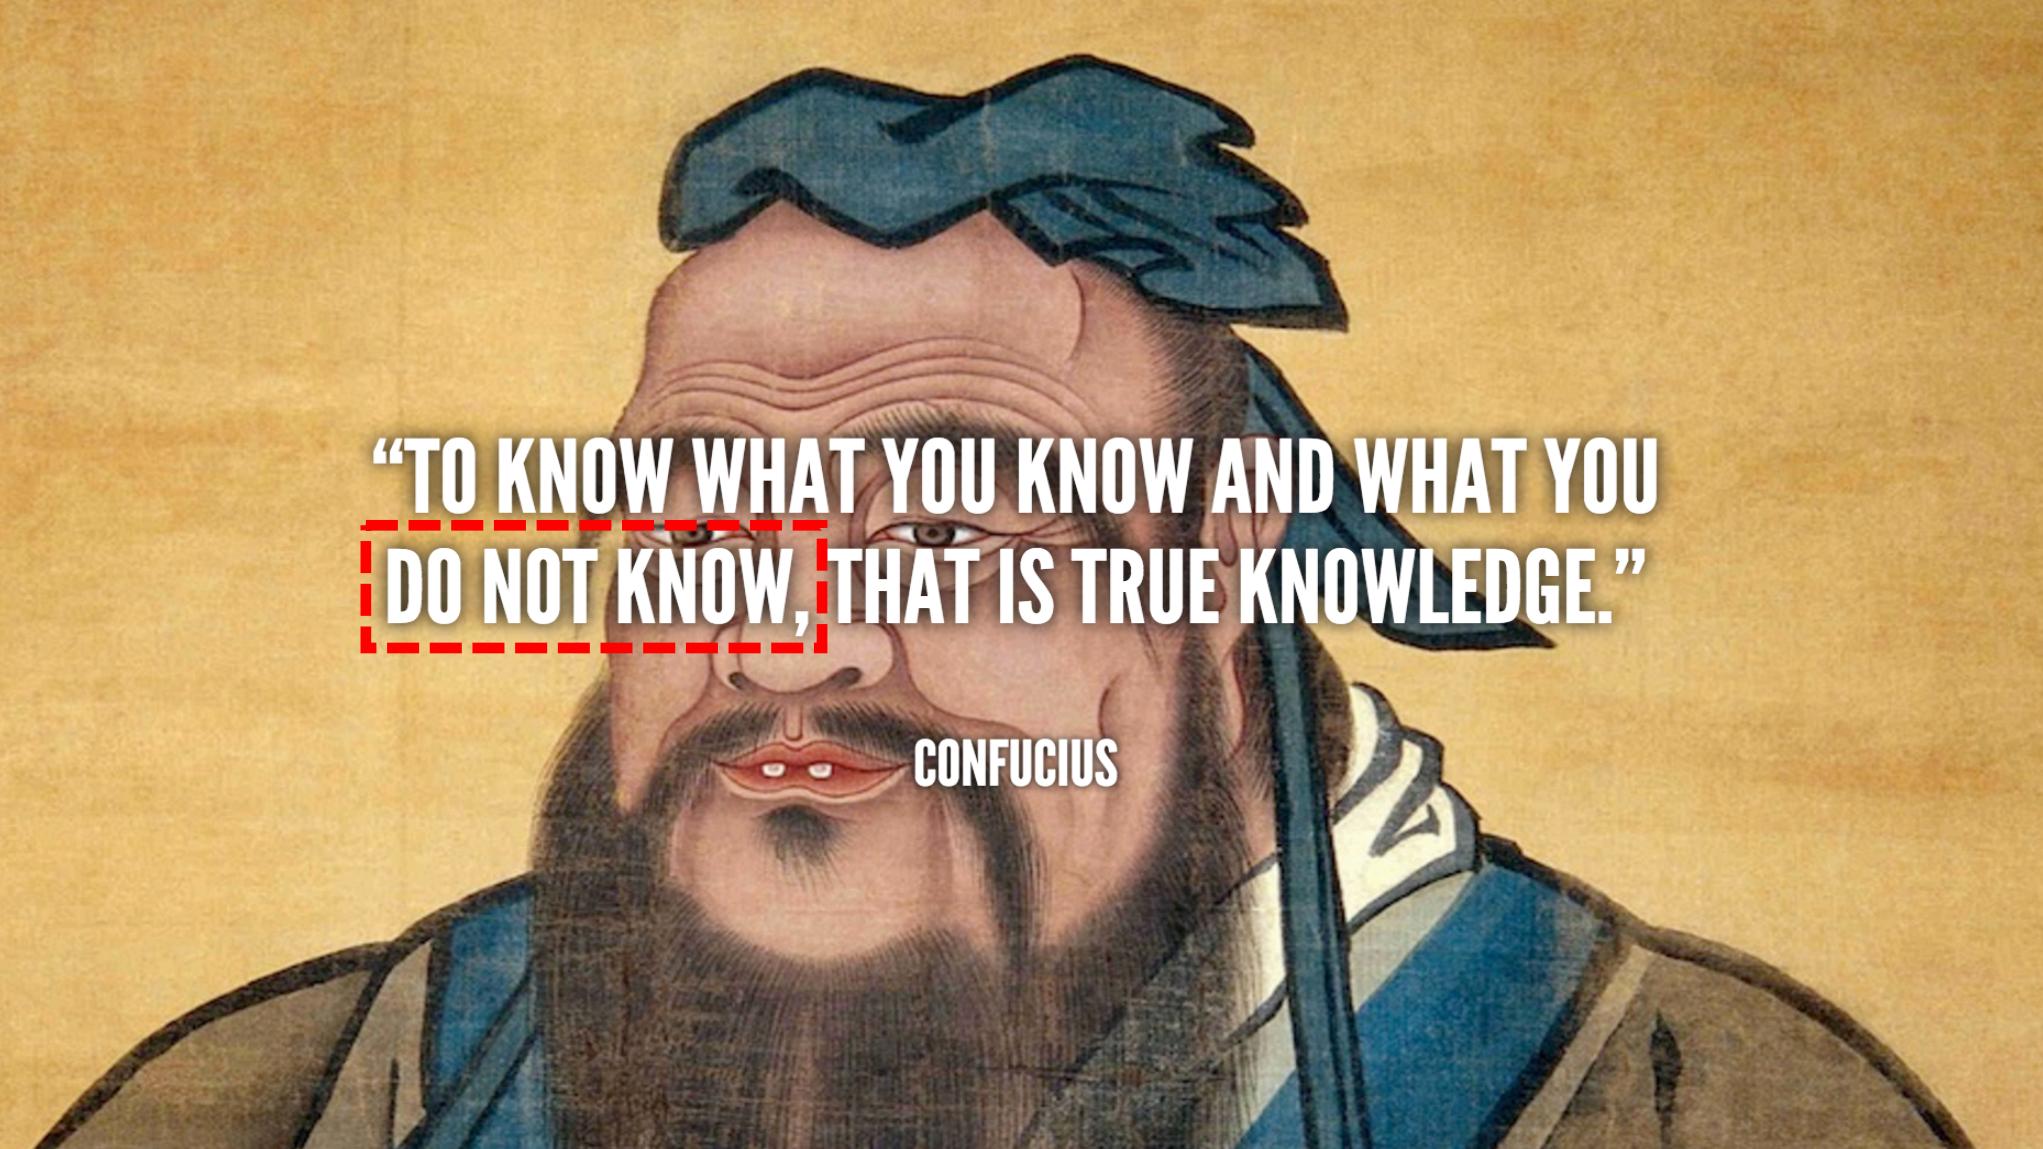


To be honest,



I will talk about **two** things



A traditional Chinese ink wash painting of a man with a long, dark beard and mustache. He has a serious expression and is wearing a blue headband and a blue and white robe. The background is a plain, light color.

**“TO KNOW WHAT YOU KNOW AND WHAT YOU
DO NOT KNOW,
THAT IS TRUE KNOWLEDGE.”**

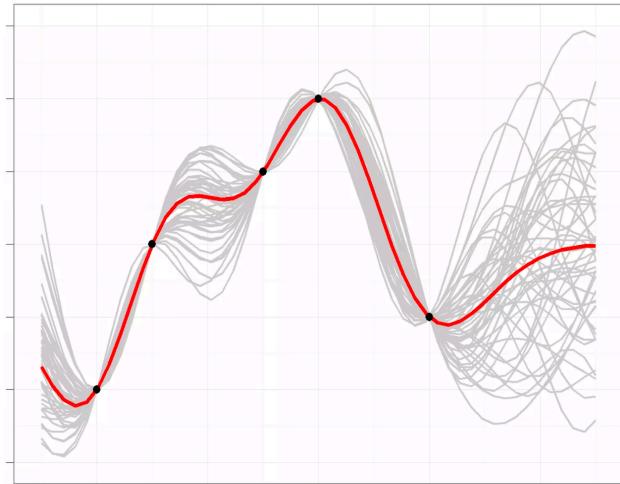
CONFUCIUS

Why?

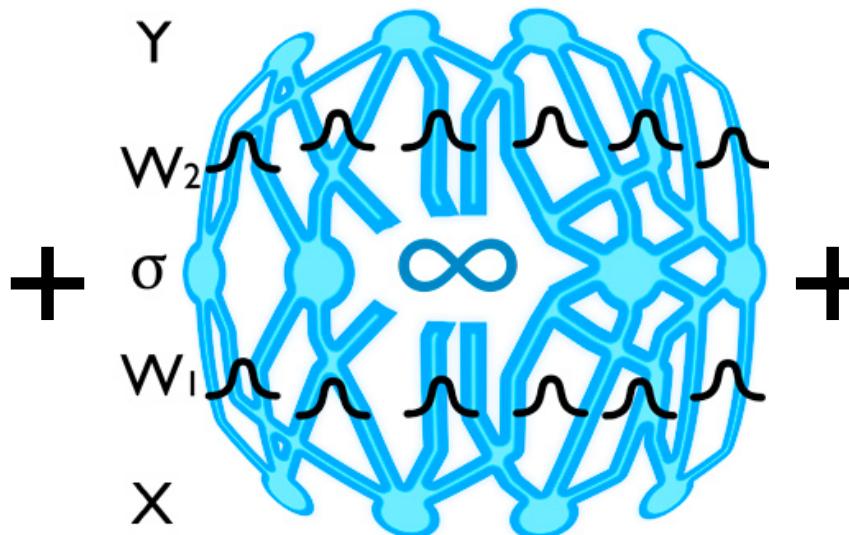


<http://www.businessinsider.com/tesla-fatal-crash-speeding-autopilot-ntsb-report-2016-7>

Then, How?



Gaussian Process

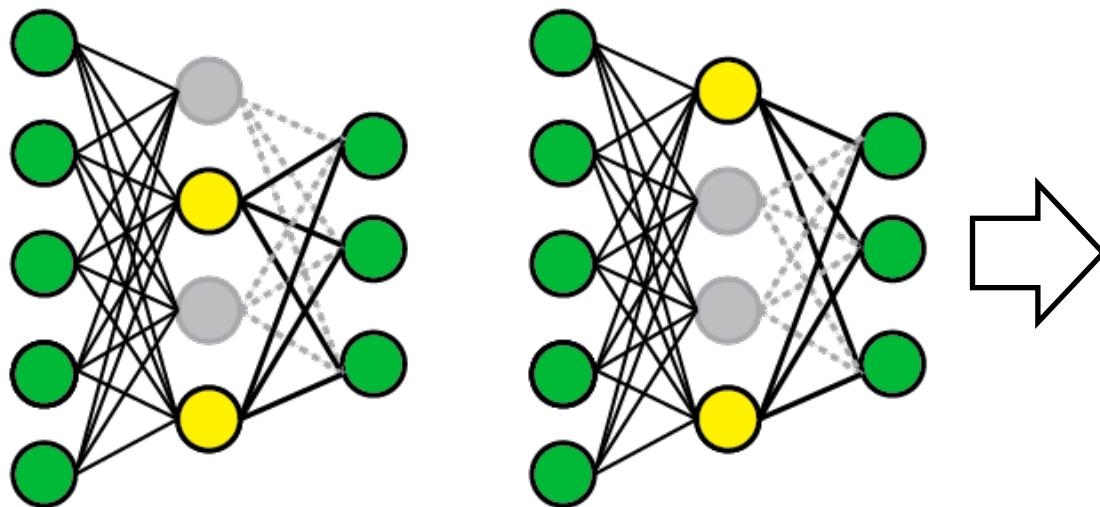


Bayesian Network

$$\begin{aligned}
 p(\mathbf{y}^* | \mathbf{x}^*, \mathbf{X}, \mathbf{Y}) &= \int p(\mathbf{y}^* | \mathbf{x}^*, \omega) p(\omega | \mathbf{X}, \mathbf{Y}) d\omega. \\
 \operatorname{argmin}_{\theta} \text{KL}(q_{\theta}(\omega) | p(\omega | \mathbf{X}, \mathbf{Y})) &= \int p(\mathbf{y}^* | \mathbf{x}^*, \omega) q_{\theta}(\omega) d\omega \\
 q_{\theta}(\mathbf{y}^* | \mathbf{x}^*) &\approx \frac{1}{T} \sum_{t=1}^T p(\mathbf{y}^* | \mathbf{x}^*, \omega_t) \\
 \mathcal{L}_{\text{VI}} &:= \int q_{\theta}(\omega) \log p(\mathbf{Y} | \mathbf{X}, \omega) d\omega - \text{KL}(q_{\theta}(\omega) || p(\omega)) \\
 \mathbf{K}(\mathbf{x}, \mathbf{y}) &= \int \mathcal{N}(\mathbf{w}; 0, l^{-2} \mathbf{I}_Q) p(b) \sigma(\mathbf{w}^T \mathbf{x} + b) \sigma(\mathbf{w}^T \mathbf{y} + b) d\mathbf{w} db \\
 \widehat{\mathbf{K}}(\mathbf{x}, \mathbf{y}) &= \frac{1}{K} \sum_{k=1}^K \sigma(\mathbf{w}_k^T \mathbf{x} + b_k) \sigma(\mathbf{w}_k^T \mathbf{y} + b_k) \\
 \mathbf{w}_k &\sim \mathcal{N}(0, l^{-2} \mathbf{I}_Q), \quad \mathbf{w}_d \sim \mathcal{N}(0, l^{-2} \mathbf{I}_K), \quad b_k \sim p(b), \\
 \mathbf{W}_1 &= [\mathbf{w}_k]_{k=1}^K, \quad \mathbf{W}_2 = [\mathbf{w}_d]_{d=1}^D, \quad \mathbf{b} = [b_k]_{k=1}^K, \\
 \omega &= \{\mathbf{W}_1, \mathbf{W}_2, \mathbf{b}\} \\
 p(\mathbf{y}^* | \mathbf{x}^*, \omega) &= \mathcal{N}\left(\mathbf{y}^*; \sqrt{\frac{1}{K}} \sigma(\mathbf{x}^* \mathbf{W}_1 + \mathbf{b}) \mathbf{W}_2, \tau^{-1} \mathbf{I}_N\right) \\
 p(\mathbf{y}^* | \mathbf{x}^*, \mathbf{X}, \mathbf{Y}) &= \int p(\mathbf{y}^* | \mathbf{x}^*, \omega) p(\omega | \mathbf{X}, \mathbf{Y}) d\omega.
 \end{aligned}$$

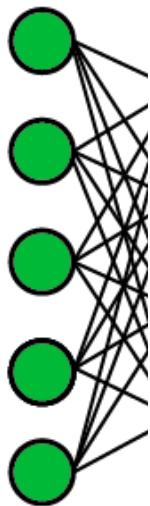
Equations

Results are pretty simple!



$$\begin{aligned}\mathbb{E}(\mathbf{y}^*) &\approx \frac{1}{T} \sum_{t=1}^T \hat{\mathbf{y}}_t^*(\mathbf{x}^*) \\ \text{Var}(\mathbf{y}^*) &\approx \tau^{-1} \mathbf{I}_D \\ &+ \frac{1}{T} \sum_{t=1}^T \hat{\mathbf{y}}_t^*(\mathbf{x}^*)^T \hat{\mathbf{y}}_t^*(\mathbf{x}^*) \\ &- \mathbb{E}(\mathbf{y}^*)^T \mathbb{E}(\mathbf{y}^*)\end{aligned}$$

Results are pretty simple!

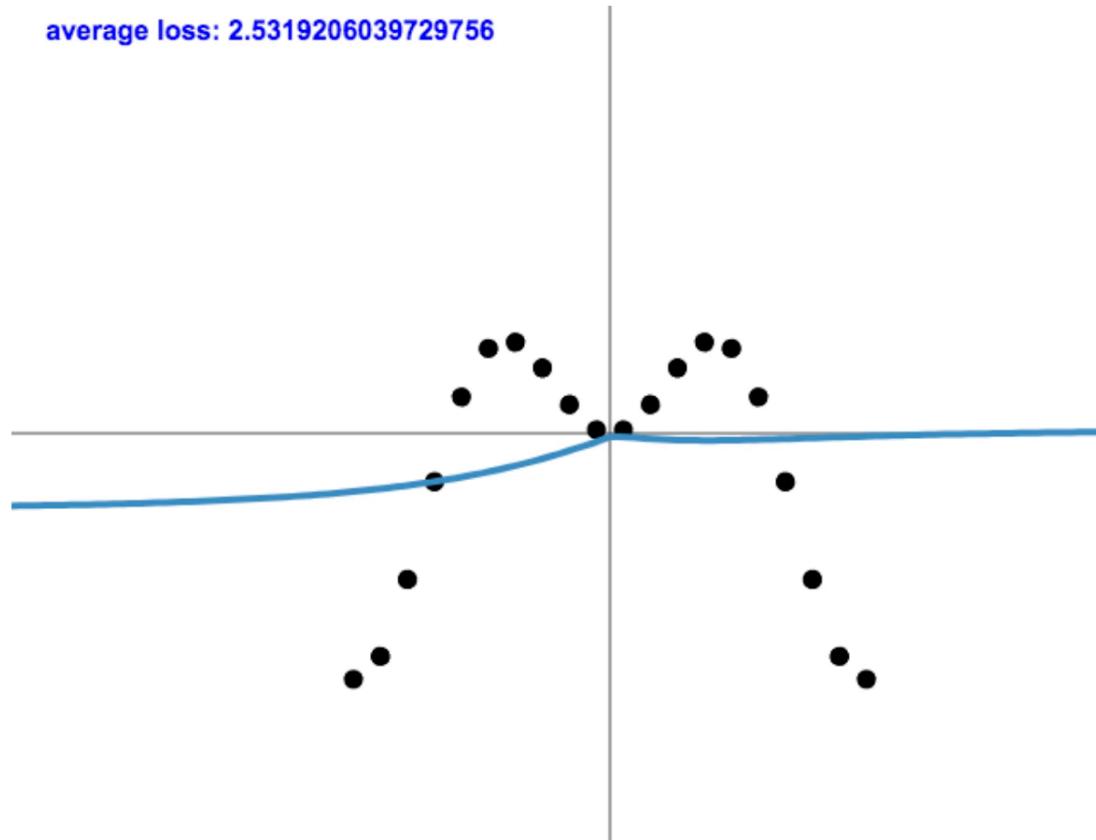


```
probs = []
for _ in xrange(T):
    probs += [model.output_probs(input_x)]
predictive_mean = numpy.mean(prob, axis=0)
predictive_variance = numpy.var(prob, axis=0)
tau = 1**2 * (1 - model.p) / (2 * N * model.weight_decay)
predictive_variance += tau**-1
```

(\mathbf{x}^*)

Cool, isn't it?

average loss: 2.5319206039729756



Uncertainty in Deep Learning



Yarin Gal

Department of Engineering
University of Cambridge

This dissertation is submitted for the degree of
Doctor of Philosophy

Gonville and Caius College

September 2016

Yarin Gal, "Uncertainty in Deep Learning", PhD Thesis, 2016

Recently,



**Deep Learning Is Not Good Enough,
We Need Bayesian Deep Learning for Safe AI**

Bayesian Deep Learning, Computer Vision, Uncertainty

http://alexgkendall.com/computer_vision/bayesian_deep_learning_for_safe_ai/

Two Types of Uncertainty

Aleatoric uncertainty

Captures our uncertainty with respect to information which our **data cannot explain**.

Aleatoric uncertainty captures noise sources such as measurement noise—noises which cannot be explained away even if more data were available.



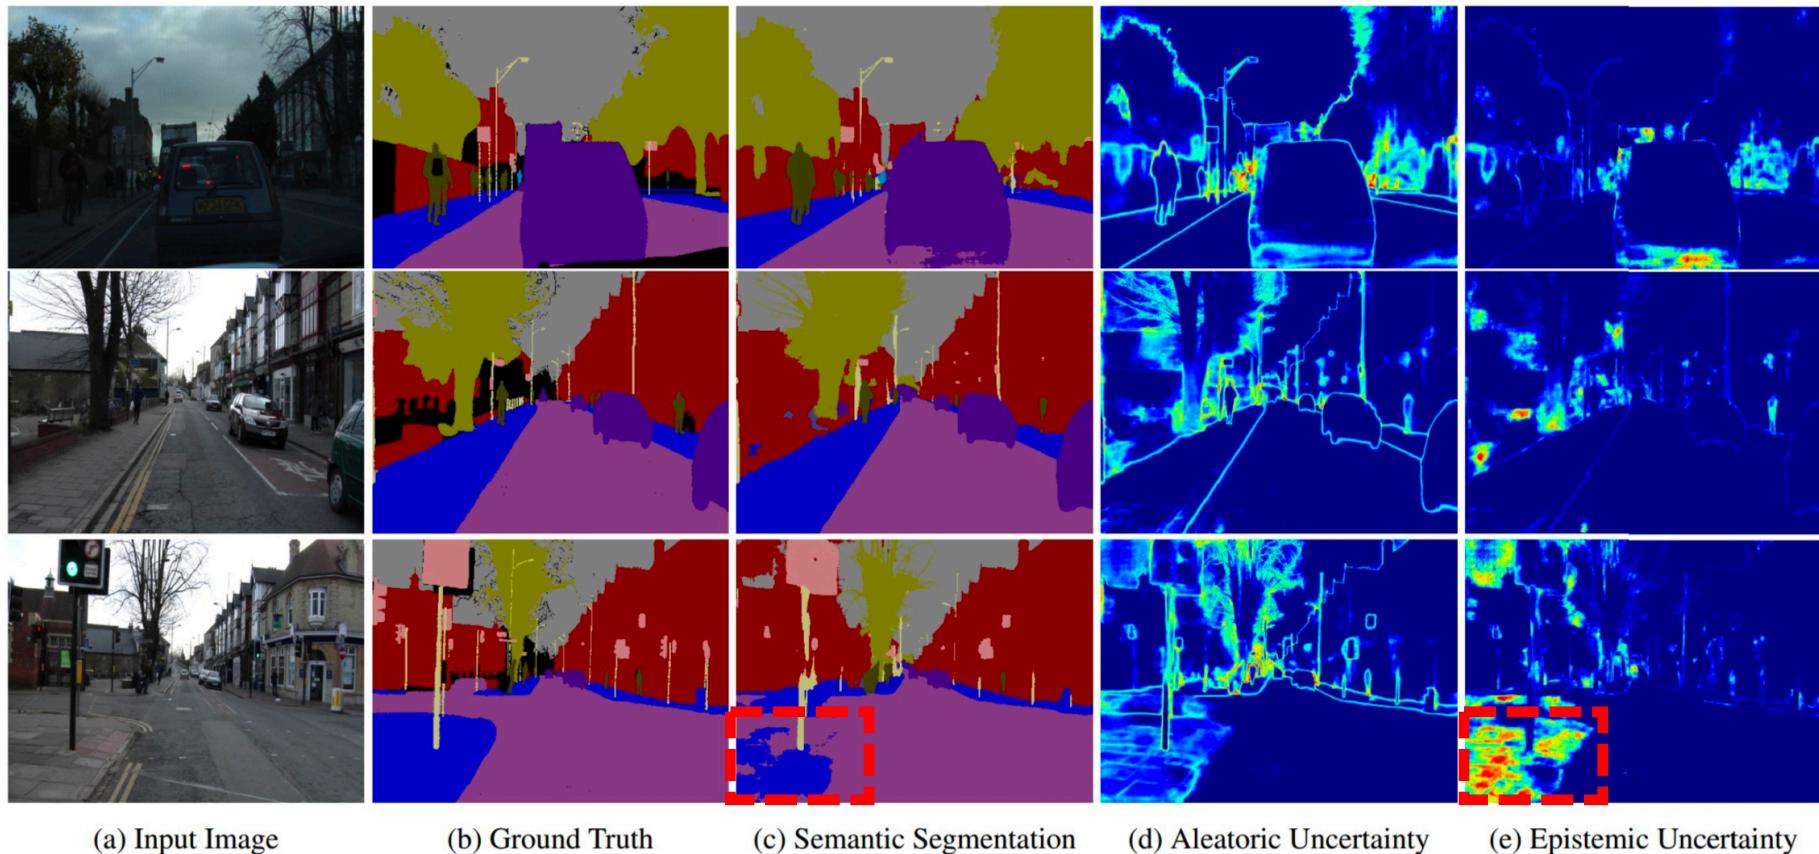
Epistemic uncertainty

Captures our **ignorance** about which model generated our collected data.

Epistemic uncertainty reduces as the amount of observed data increases—hence its alternative name “reducible uncertainty”.



Semantic Segmentation Results



(a) Input Image

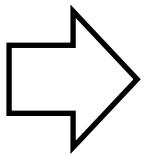
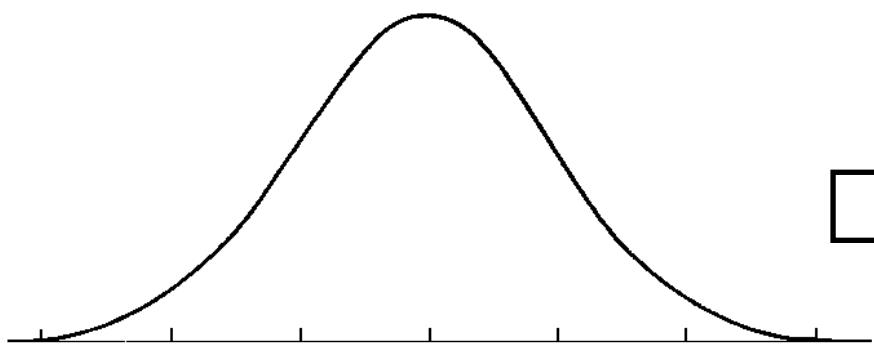
(b) Ground Truth

(c) Semantic Segmentation

(d) Aleatoric Uncertainty

(e) Epistemic Uncertainty

How They Did It?



$$\text{Var}(x) \approx \frac{1}{T} \sum_{t=1}^T \hat{x}_t^2 - \left(\frac{1}{T} \sum_{t=1}^T \hat{x}_t \right)^2 + \frac{1}{T} \sum_{t=1}^T \hat{\sigma}_t^2$$

?

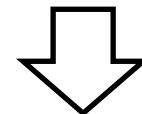
$$[\hat{x}, \hat{\sigma}^2] = f^{\widehat{\mathbf{W}}}(I)$$

Density Network

Law of Total Variance

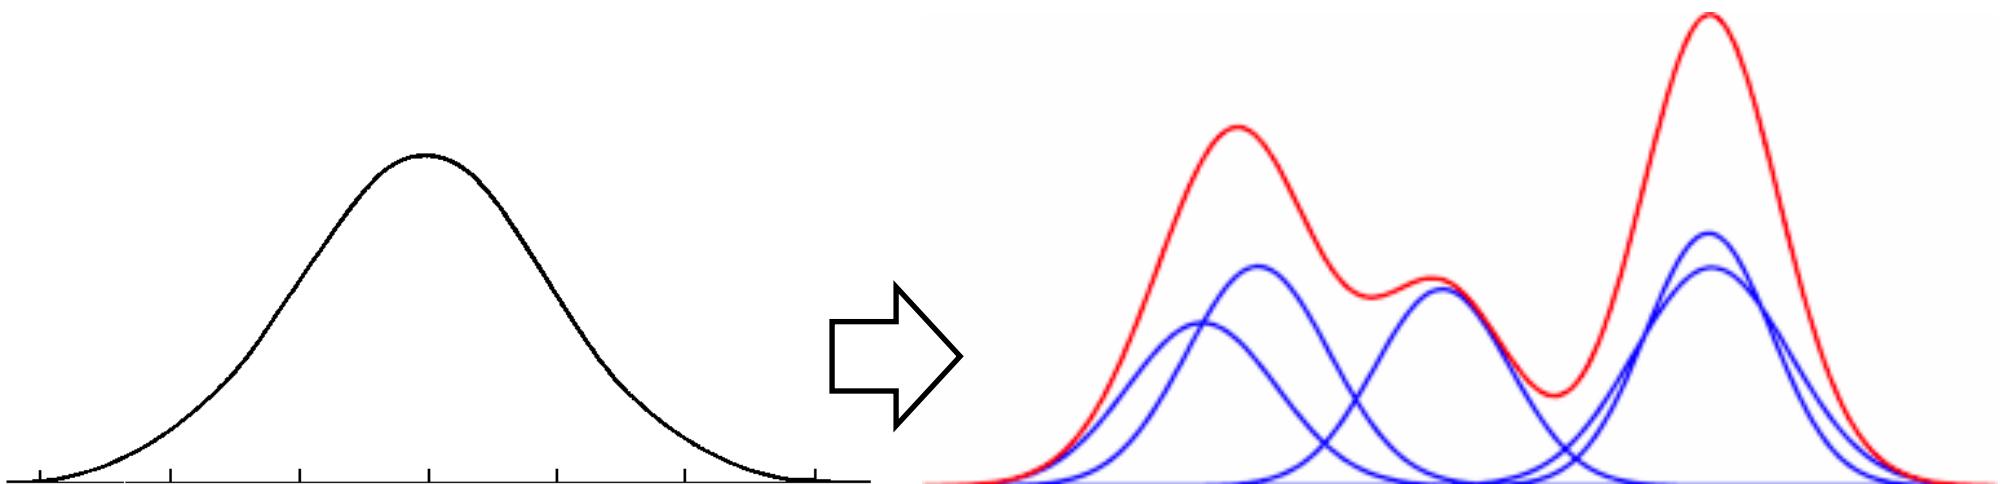
$$V[X] = \boxed{V(E[X|y])} + \boxed{E(V[X|y])}$$

Explained Variance Unexplained Variance



$$\text{Var}(x) \approx \boxed{\frac{1}{T} \sum_{t=1}^T \hat{x}_t^2 - \left(\frac{1}{T} \sum_{t=1}^T \hat{x}_t \right)^2} + \boxed{\frac{1}{T} \sum_{t=1}^T \hat{\sigma}_t^2}$$

What We Did?



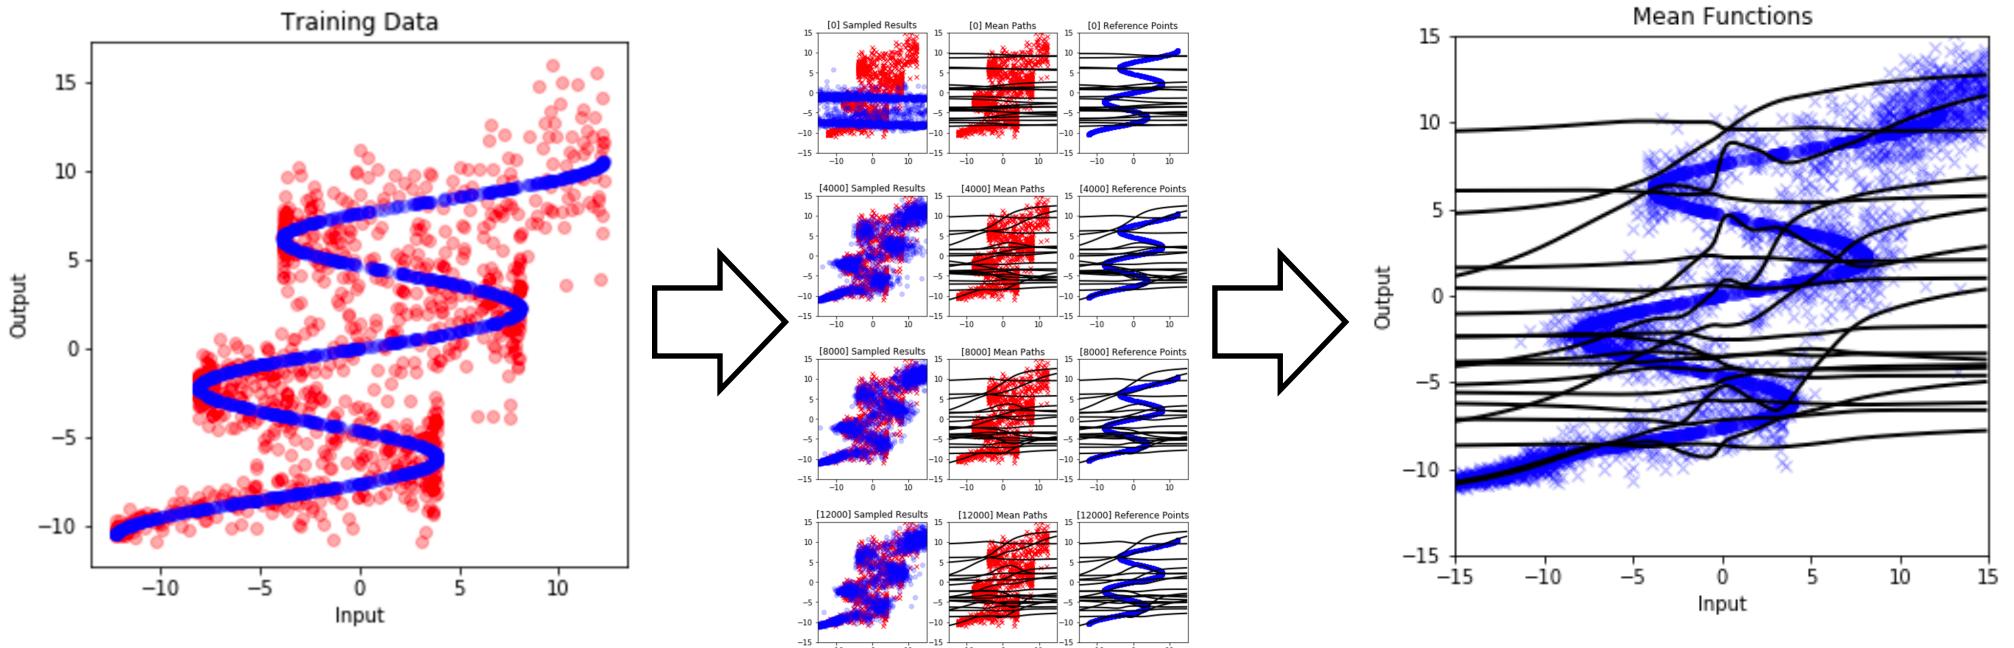
$$[\hat{x}, \hat{\sigma}^2] = f^{\widehat{\mathbf{W}}}(I)$$

Density Network

$$p(\mathbf{x}) = \sum_{k=1}^K \pi_k N(\mathbf{x} | \mu_k, \Sigma_k)$$

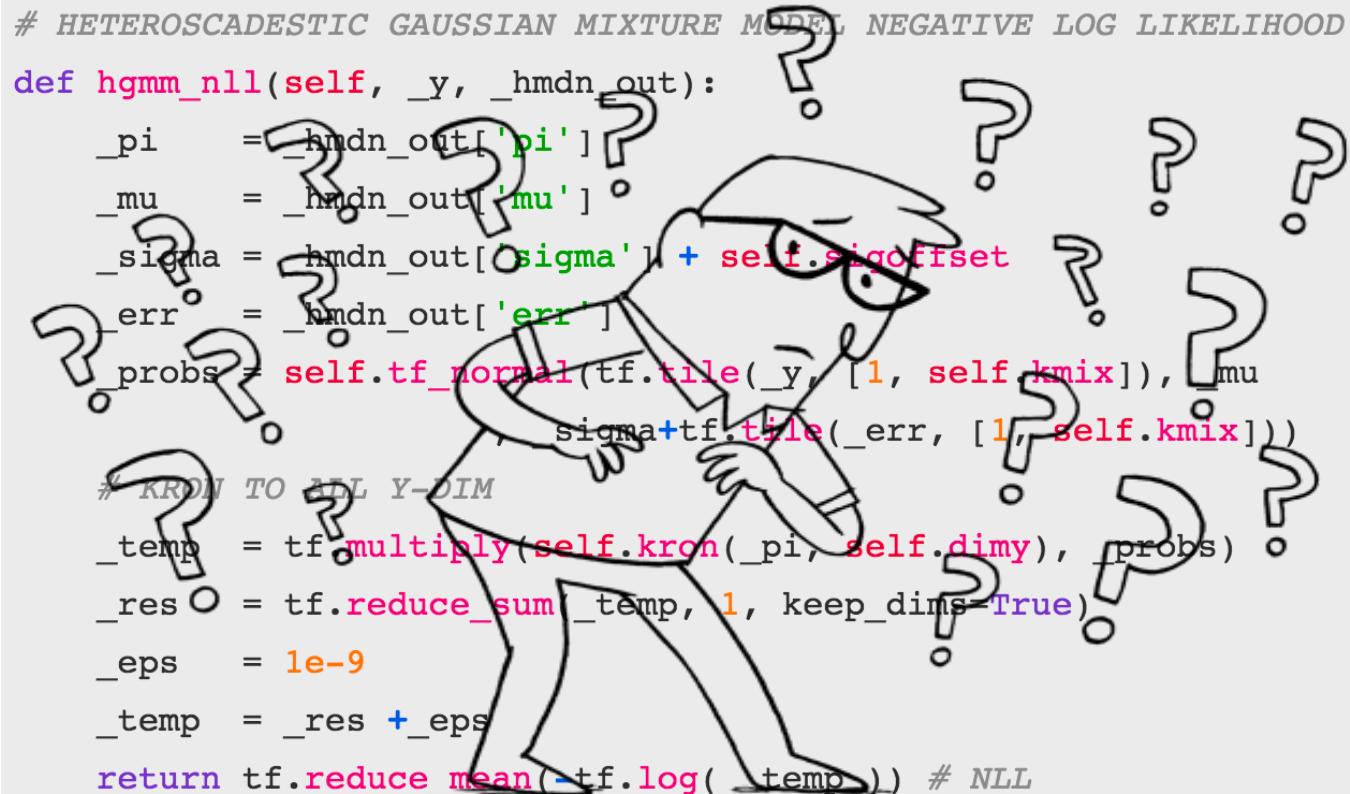
Mixture Density Network

Mixture Density Network



<https://github.com/sjchoi86/advanced-tensorflow/blob/master/mdn/hmdn.ipynb>

Mixture Density Network

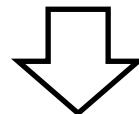


```
# HETEROSCADESTIC GAUSSIAN MIXTURE MODEL NEGATIVE LOG LIKELIHOOD
def hgmm_nll(self, _y, _hmdn_out):
    _pi      = _hmdn_out['pi']
    _mu     = _hmdn_out['mu']
    _sigma  = _hmdn_out['sigma'] + self.sigoffset
    _err    = _hmdn_out['err']
    _probs = self.tf_normal(tf.tile(_y, [1, self.kmix]), _mu
                           , sigma+tf.tile(_err, [1, self.kmix]))
    # KRON TO ADD Y-DIM
    _temp   = tf.multiply(self.kron(_pi, self.dimy), _probs)
    _res    = tf.reduce_sum(_temp, 1, keep_dims=True)
    _eps    = 1e-9
    _temp   = _res + _eps
    return tf.reduce_mean(-tf.log(_temp)) # NLL
```

Revisit Law of Total Variance

$$V[X] = \boxed{V(E[X|y])} + \boxed{E(V[X|y])}$$

Explained Variance Unexplained Variance

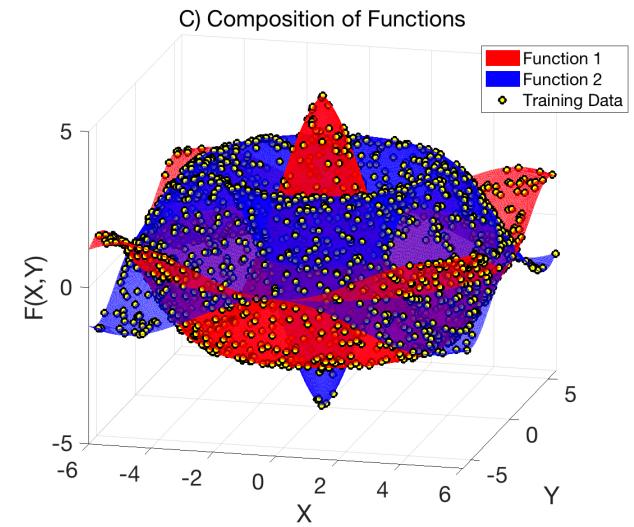
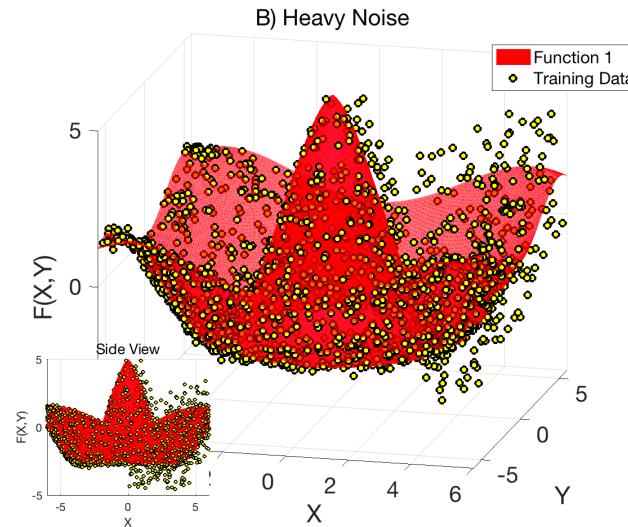
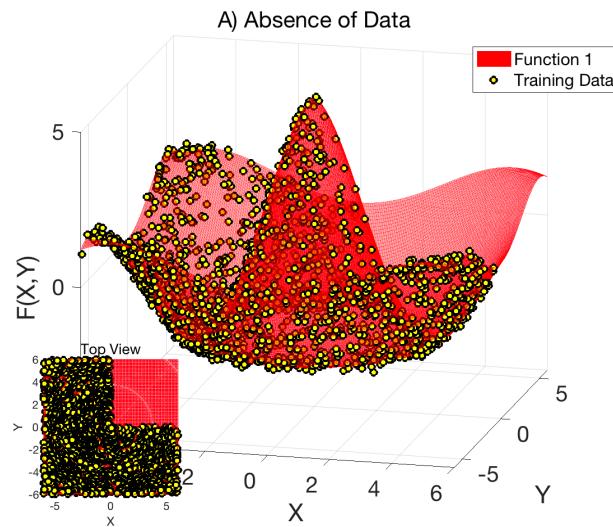


$$V[X] = \boxed{\sum_y p(y)(E[X|y] - E[X])^2} + \boxed{\sum_y p(y) V(X|y)}$$

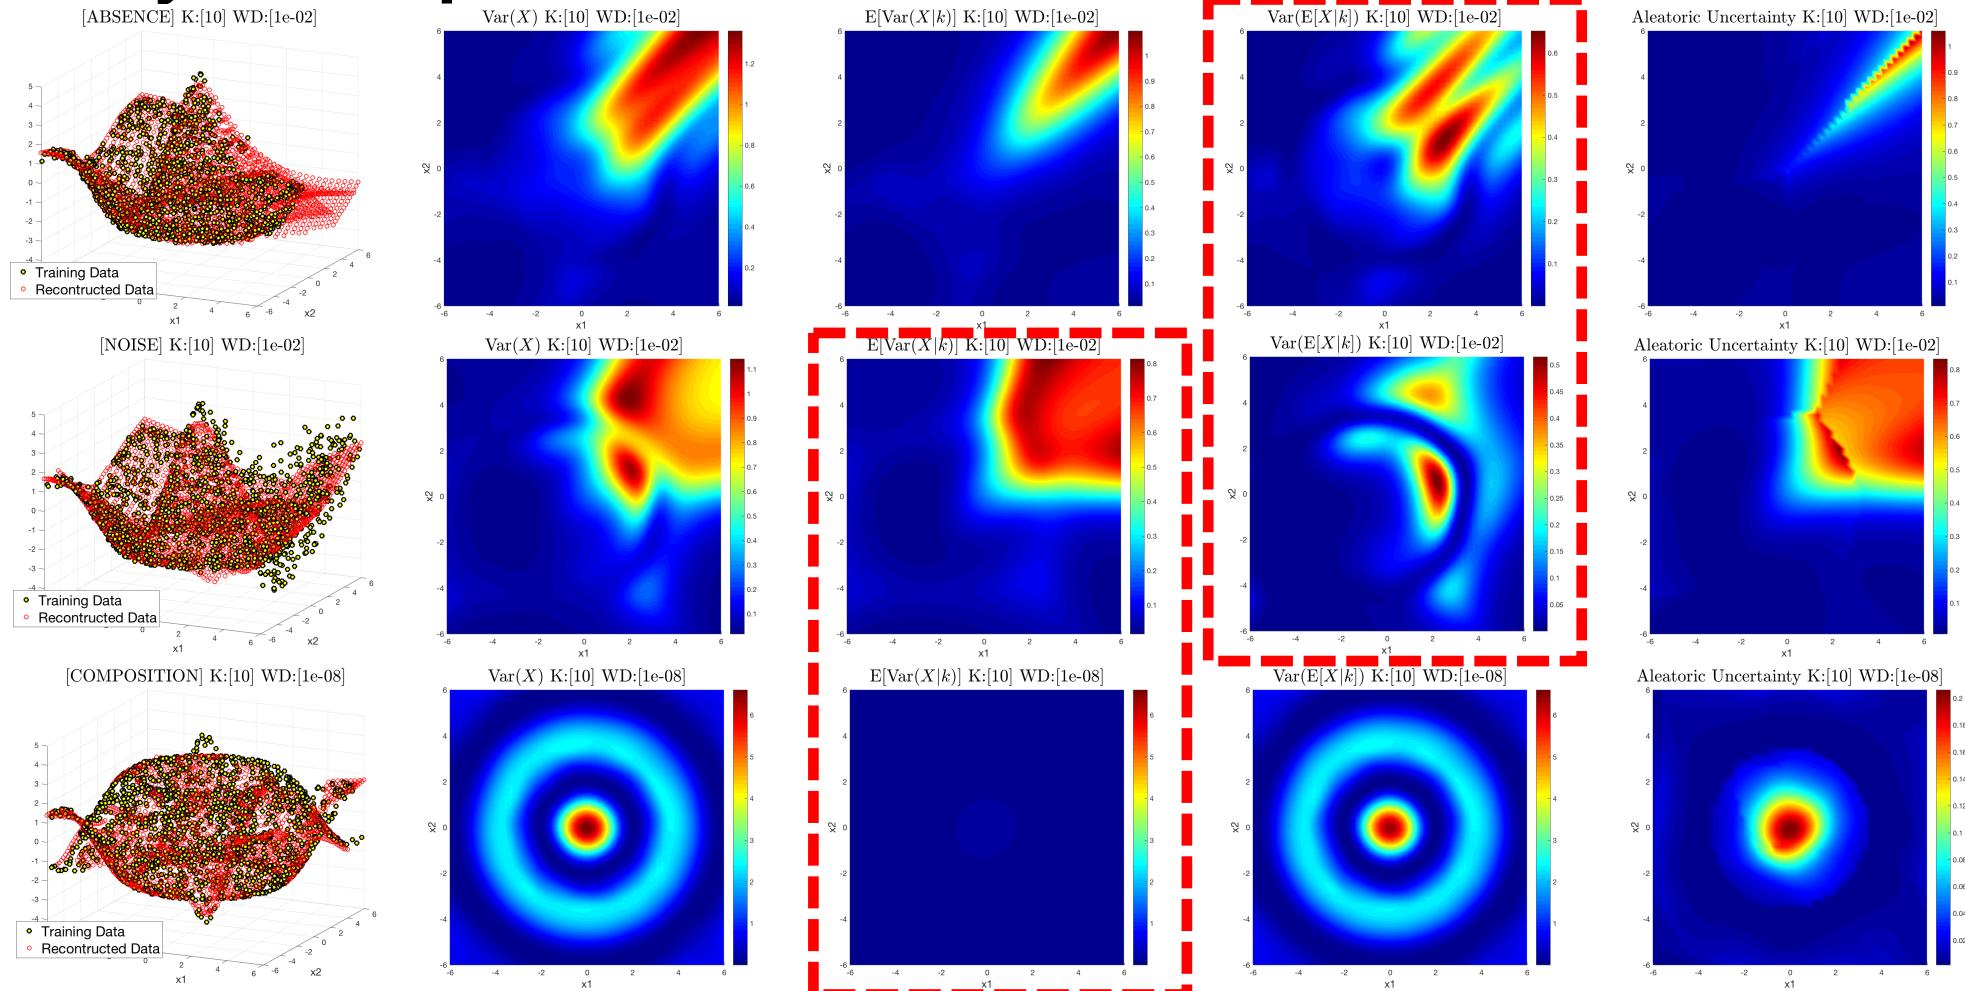
Depends on mean functions Depends on variance functions

Toy Example

Three Scenarios

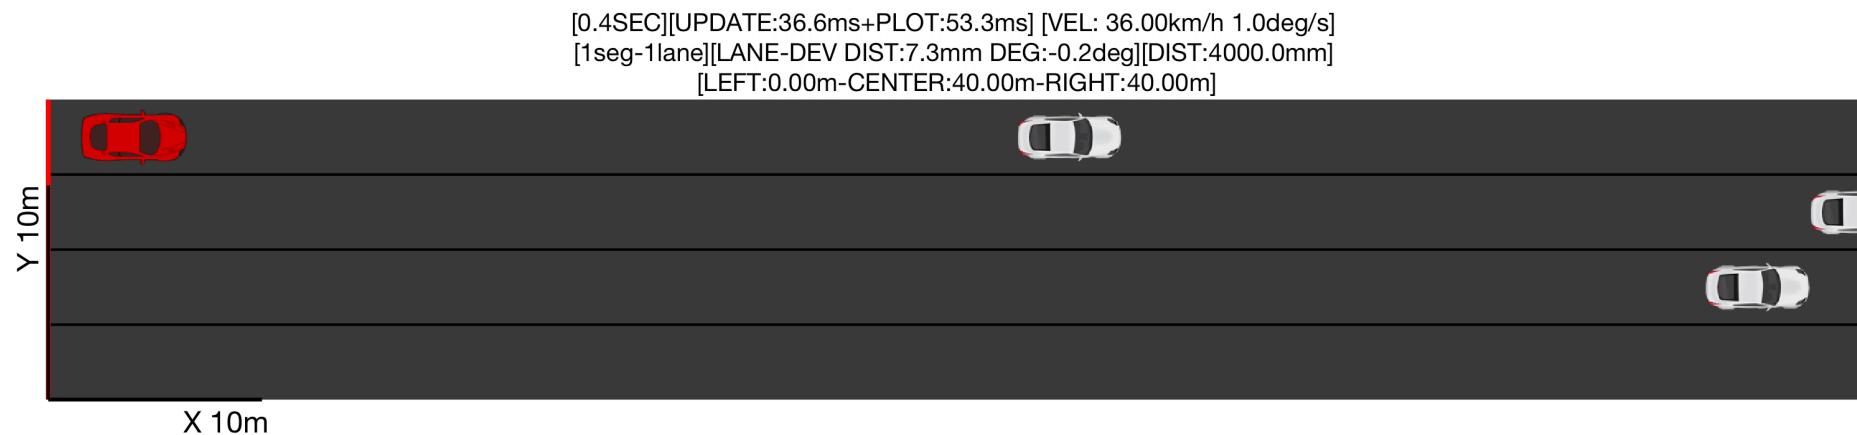


Toy Example



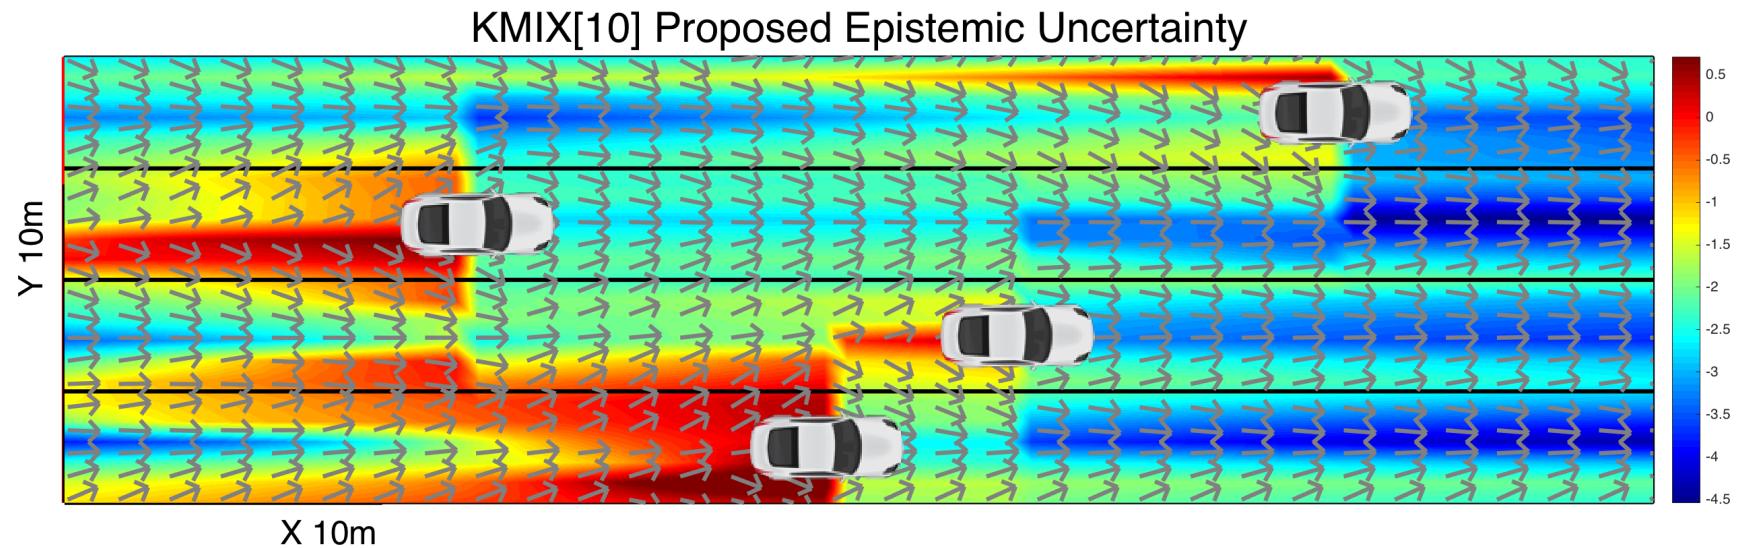
Track Driving Scenario

Demonstration Collection



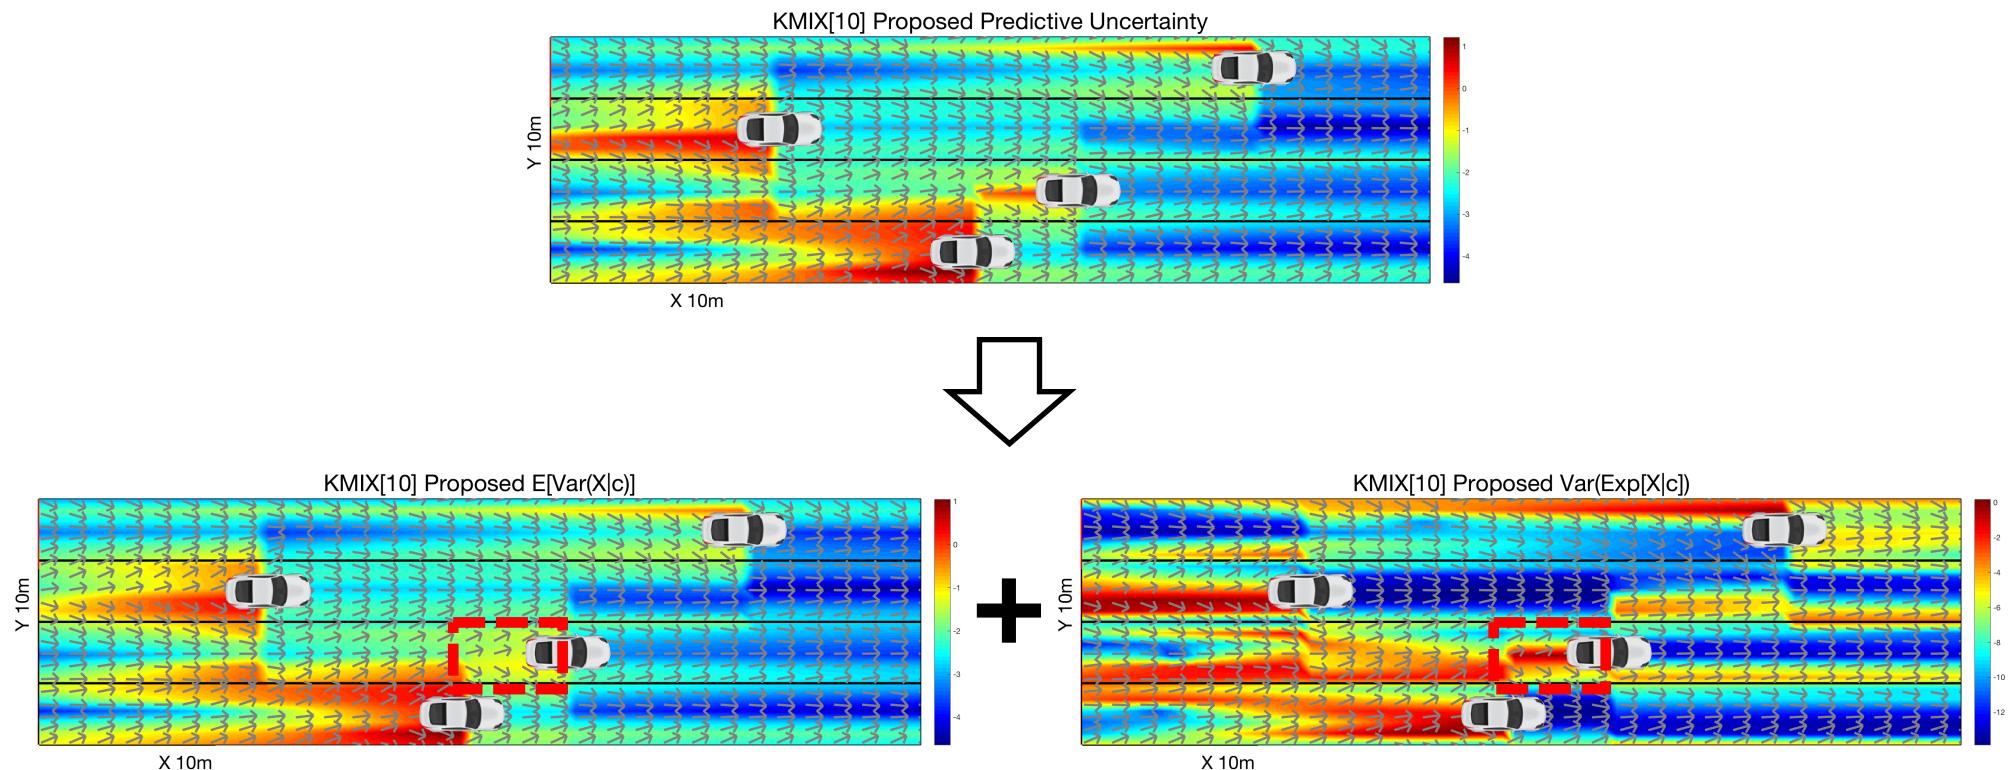
Track Driving Scenario

Predictive Uncertainty



Track Driving Scenario

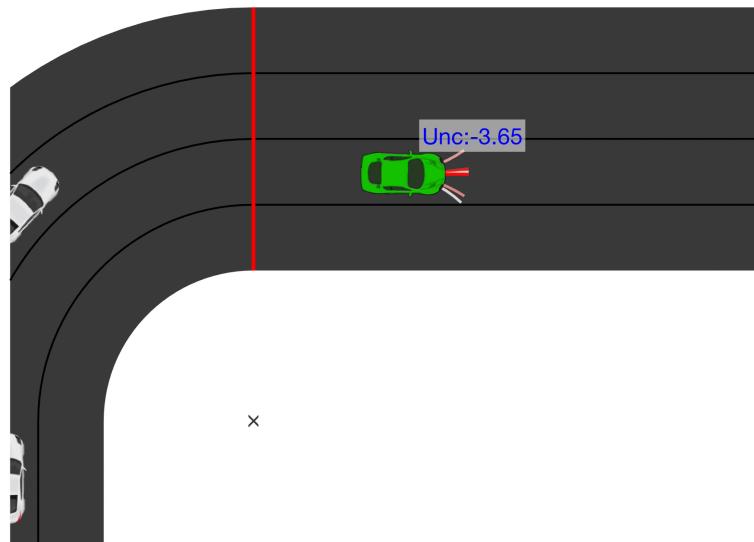
Uncertainty Decompositions



Track Driving Scenario

Track Simulator

Naïve Navigation



Uncertainty **below -1**
Directional vel: 20m/s

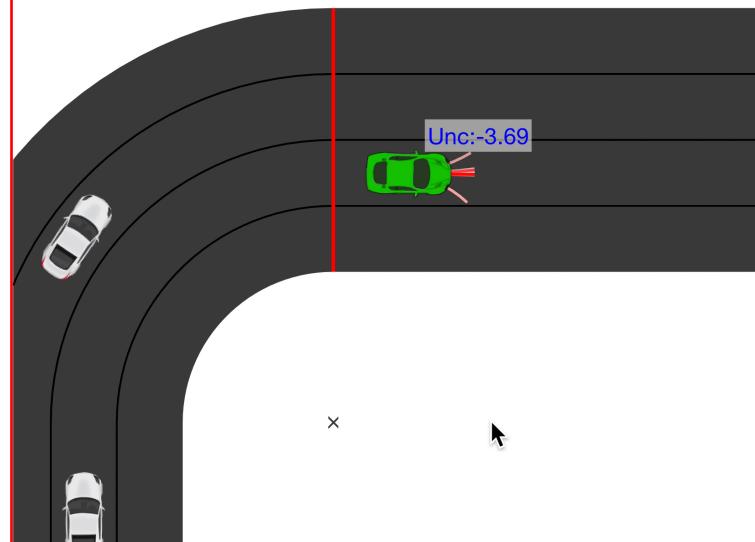


Uncertainty **-1~0**
Directional vel: 15m/s



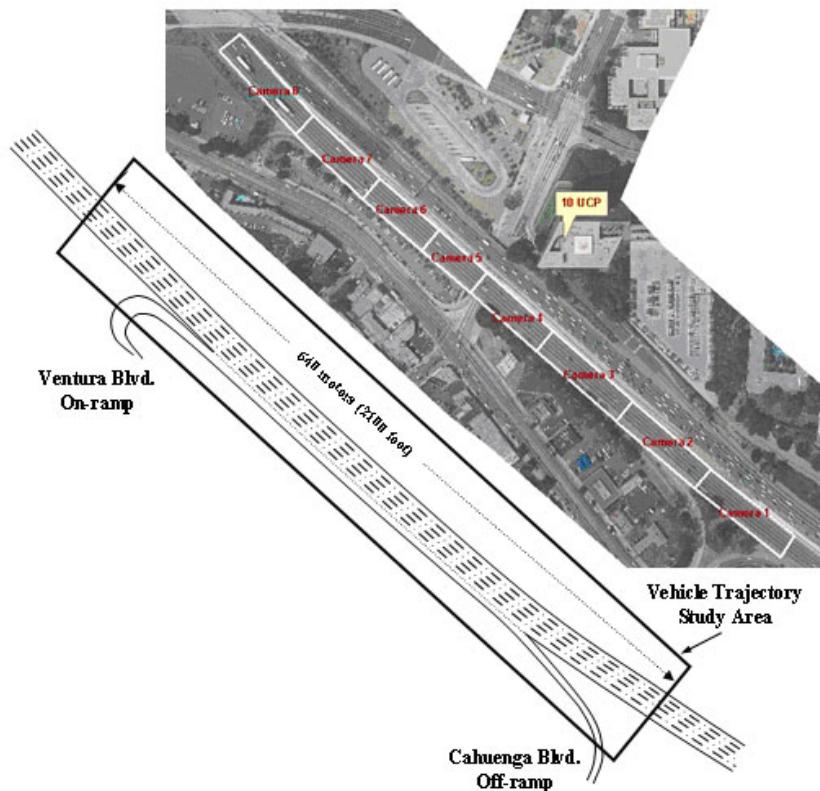
Uncertainty **above 0**
Directional vel: 10m/s

Uncertainty-aware Navigation



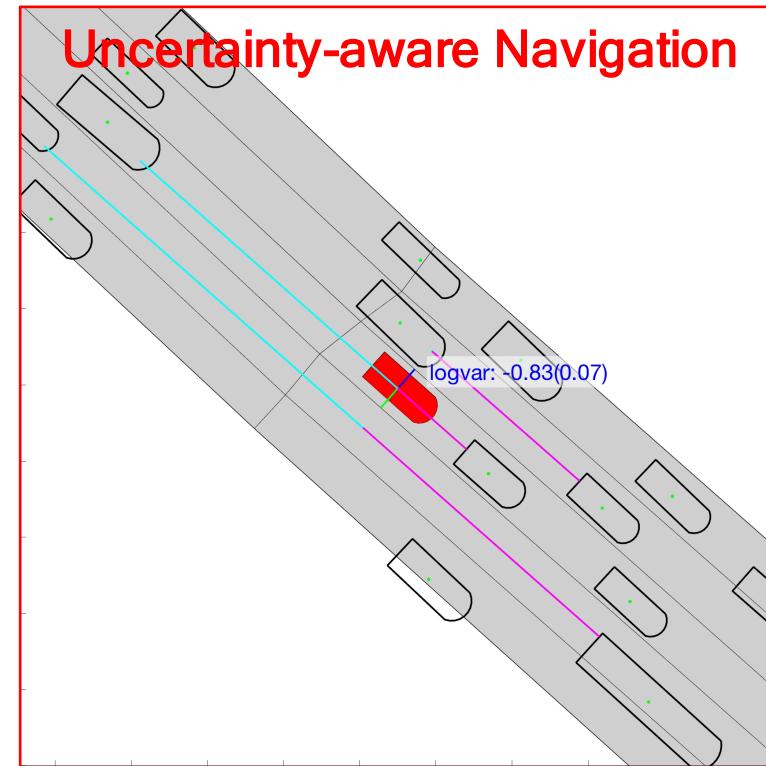
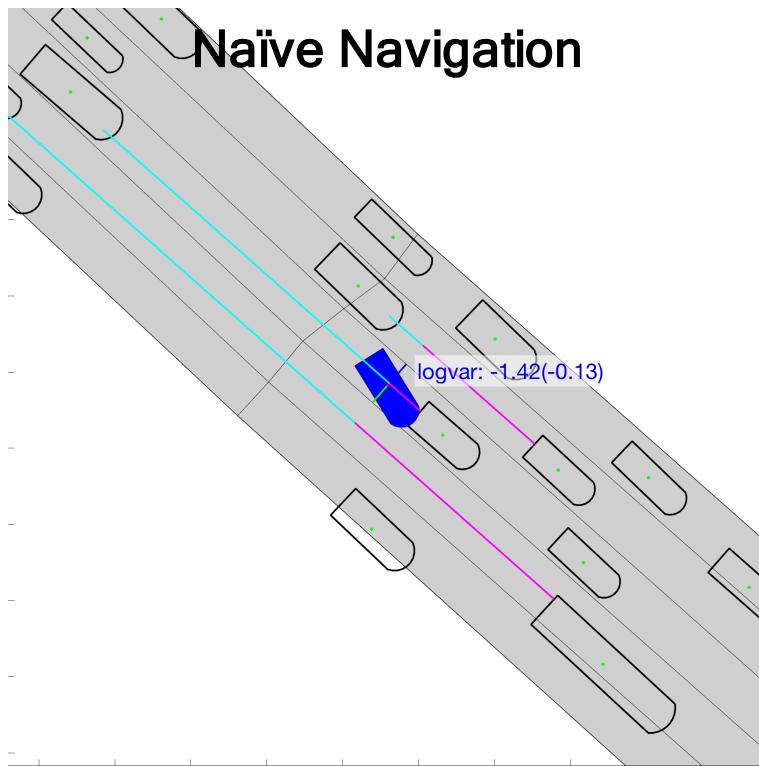
Track Driving Scenario

US Highway 101



Track Driving Scenario

US Highway 101



Conclusion

A simple but powerful uncertainty estimation method is presented with a mixture density network.

Successfully applied to uncertainty-aware learning from demonstration.

Currently, writing a paper about this topic. ☺

**Uncertainty-Aware Learning from Demonstration
with Mixture Density Networks**

Sungjoon Choi, Kyungjae Lee, Sungbin Lim, and Songhwai Oh

Special thanks to Sungbin Lee and Kyungjae Lee.

